

Contents

1	Reading data from files	2
1.1	The working directory (<code>getwd()</code> and <code>setwd()</code>)	2
2	Exercise 2.1	2
2.1	Importing data	2
2.2	Sorting a data frame (<code>order()</code>)	4
3	Filtering entries	5
4	Exercise 2.2	6
4.1	Importing fixed width format files <code>read.fwf()</code>	6
5	Exercise 2.3	8
6	Exercise 2.4	9
7	Exercise 2.5	9
8	Exercise 2.6	11
9	Exercise 2.7	11

1 Reading data from files

For the most part, in this course, we will either get data from other packages in R, or read data from files. It is important to be familiar with how to import data, and generally there are only two (three) functions to be aware of, `read.fwf()` and `read.table()` (and `read.csv()`, which is really just `read.table()` with some preset settings).

In general, the first step is to be aware what is the structure of the data. Open the data file you want to import and study its structure. Does it have a header? Are the individual data points in an observation separated by anything? And in general, most files (almost all) should have one observation in one row.

1.1 The working directory (`getwd()` and `setwd()`)

If you encounter any file-related errors (e.g. file not found) when trying to read a file, the first thing you should check is your working directory. Most problems will likely be solved if you check your working directory, because **any functions that require a file (filepath) will try to find the file (filepath) relative to your current working directory.**

To see your current working directory, use `getwd()`. You may, for example, get the following response: `C:/Users/caleb/Documents`. Thus, if I run `read.table("lab1fixed.txt")`, it will look for the file at `C:/Users/caleb/Documents/lab1fixed.txt`. So, check if the file is where you intend it to be, or check that your working directory is where you intend it to be.

To change your working directory, you can use `setwd()` and pass it a path.

Alternatively, in RStudio, use the Files tab, click the 3 dots at the right of the panel and navigate to the folder with your files. Click the cog icon, and find the 'Set As Working Directory' option.

Both methods are equivalent, as they ultimately use the `setwd()` function - it will appear in the R console.


A quick mention for Rmarkdown (.Rmd) files: I find that as long as the files are in the same folder (e.g. `Documents/script.Rmd` and `Documents/data.csv`) the working directory is irrelevant.

2 Exercise 2.1

2.1 Importing data

The lab sheet suggests using `read.table()`. To understand how to use it, look at the documentation with `?read.table`. The usage indicates there are many parameters the function can accept, but the first few are the most important: `file`, `header`, and `sep`.

`file` is simply the filename of the data we want to import.



header indicates whether the first line of the file contains information on the column headers. In `read.table()`, the default value is FALSE. Looking at the data file, we can see that the first line indeed are labels for each column, so we must change header to TRUE.

`sep` defines the character that separates each column. The default value is an empty string, which will separate by white space, i.e. any amount of space characters, tab characters, line breaks. In this file, it is a comma (,), so we change `sep` to ', '.

```
car1 <- read.table("Lab02_Ex1.csv", header=T, sep=",")  
car1
```

```
## # A tibble: 12 x 7  
##   Car           MPG Displacement Horsepower Weight Model Origin  
##   <chr>        <dbl>      <int>       <int>    <int> <int> <chr>  
## 1 Peugeot 304     30          79        70    2074    71 Europe  
## 2 Audi 100LS      20         114        91    2582    73 Europe  
## 3 Volkswagen Rabbit  29          90        70    1937    75 Europe  
## 4 Volvo 245       20         130       102    3150    76 Europe  
## 5 Toyota Corolla    25         113        95    2228    71 Japan  
## 6 Subaru FT        26         108        93    2391    74 Japan  
## 7 Honda Civic CVCC 36.1         91        60    1800    78 Japan  
## 8 Datsun 810       22         146        97    2815    77 Japan  
## 9 Ford F250         10         360       215    4615    70 US
```

```
## 10 Chevy C20          10        307      200    4376    70 US
## 11 Pontiac Catalina   16        400      170    4668    75 US
## 12 Chevrolet Chevette 30        98       68     2155    78 US
```

Alternatively, one could have used the `read.csv()` function directly, since it has different default values which you can see with `?read.csv`.

```
car1 <- read.csv("Lab02_Ex1.csv")
car1
```

```
## # A tibble: 12 x 7
##   Car           MPG Displacement Horsepower Weight Model Origin
##   <chr>        <dbl>      <int>        <int>  <int> <chr>
## 1 Peugeot 304  30         79          70    2074  71 Europe
## 2 Audi 100LS   20         114         91    2582  73 Europe
## 3 Volkswagen Rabbit 29         90          70    1937  75 Europe
## 4 Volvo 245    20         130         102   3150  76 Europe
## 5 Toyota Corolla 25         113         95    2228  71 Japan
## 6 Subaru FT    26         108         93    2391  74 Japan
## 7 Honda Civic CVCC 36.1        91          60    1800  78 Japan
## 8 Datsun 810   22         146         97    2815  77 Japan
## 9 Ford F250    10         360         215   4615  70 US
## 10 Chevy C20   10         307        200    4376  70 US
## 11 Pontiac Catalina 16         400        170    4668  75 US
## 12 Chevrolet Chevette 30         98         68    2155  78 US
```

2.2 Sorting a data frame (`order()`)

To sort a data frame, we can use the `order()` function, which just returns a numerical ranking to the data. For example:

```
x <- c(100, 200, 150, 175, 125)
x[c(3,3,2,4,4)] # get the 3rd, 3rd, 2nd, 4th, 4th entries.
```

```
## [1] 150 150 200 175 175
```

```
order(x) # 1, 5, 3, 4, 2
```

```
## [1] 1 5 3 4 2
```

```
x <- x[order(x)]
x
```

```
## [1] 100 125 150 175 200
```

In earlier exercises, we used a boolean condition (true or false, based on some criteria) to determine whether to keep or discard a row (observation). We can also use a numerical index to determine what entries to take, as in line 2 in the above code block.

You can see that `order()` gives the ordering of the numbers from smallest (1) to largest (5). So, when we order `x` by this sequence of indexes, we will get the elements in `x` ordered from smallest to largest.

Since we want to order the entries by weight, we use the `order()` function on the weight vector `car1$Weight`. By default, `order()` will return indexes in an increasing order. This behaviour can be controlled by changing the optional parameter `decreasing`, whose default value is `FALSE`.

```
ordering <- order(car1$Weight, decreasing=T)
car1 <- car1[ordering,]
car1
```

```
## # A tibble: 12 x 7
##   Car           MPG Displacement Horsepower Weight Model Origin
##   <chr>      <dbl>       <int>       <int>    <int> <int> <chr>
## 1 Pontiac Catalina     16          400        170    4668     75 US
## 2 Ford F250            10          360        215    4615     70 US
## 3 Chevy C20            10          307        200    4376     70 US
## 4 Volvo 245            20          130        102    3150     76 Europe
## 5 Datsun 810            22          146         97    2815     77 Japan
## 6 Audi 100LS           20          114         91    2582     73 Europe
## 7 Subaru FT             26          108         93    2391     74 Japan
## 8 Toyota Corolla        25          113         95    2228     71 Japan
## 9 Chevrolet Chevette    30           98         68    2155     78 US
## 10 Peugeot 304           30           79         70    2074     71 Europe
## 11 Volkswagen Rabbit    29           90         70    1937     75 Europe
## 12 Honda Civic CVCC     36.1        91         60    1800     78 Japan
```

3 Filtering entries

We have seen this in action last week. In part (d), we have a compound condition to filter for. To do so, we simply get the two boolean vectors representing the conditions, and do a logical AND operation on them using `&`.

Below, it is interpreted as “take the rows where weight is greater than 3000 and horsepower is greater than 200, and all the columns.”

```
car1[car1$Weight > 3000 & car1$Horsepower > 200,]

## # A tibble: 1 x 7
##   Car           MPG Displacement Horsepower Weight Model Origin
##   <chr>        <dbl>       <int>       <int>    <int> <chr>
## 1 Ford F250     10         360        215     4615    70 US
```

The `subset()` function is another way to filter entries, with similar syntax. The columns here can be referred to as variables directly, without specifying the data frame, i.e. `car1$` is not necessary. An optional parameter `select` can be specified to determine what columns to keep.

```
subset(car1, Weight > 3000 & Horsepower > 200)

## # A tibble: 1 x 7
##   Car           MPG Displacement Horsepower Weight Model Origin
##   <chr>        <dbl>       <int>       <int>    <int> <chr>
## 1 Ford F250     10         360        215     4615    70 US

subset(car1, Weight > 3000 & Horsepower > 200, select=Car:Weight)

## # A tibble: 1 x 5
##   Car           MPG Displacement Horsepower Weight
##   <chr>        <dbl>       <int>       <int>    <int>
## 1 Ford F250     10         360        215     4615
```

4 Exercise 2.2

4.1 Importing fixed width format files `read.fwf()`

Open and observe the format of `Lab02fixed.txt`. As the name suggests, it is a **fixed-width format** file, with no separators between each data point in a single observation. Instead, some number of columns corresponds to some information about a single observation, with the details for this file provided in the lab sheet. These columns are passed in under the `widths` parameter.

```
lab2 <- read.fwf("Lab02fixed.txt", widths = c(3,1,3,2,1),
                  col.names = c("id", "gender", "height", "weight",
                               "siblings"))

str(lab2)
```

```
## 'data.frame': 120 obs. of 5 variables:
## $ id      : int 201 202 203 204 205 206 207 208 209 210 ...
## $ gender   : chr "M" "F" "F" "F" ...
## $ height   : int 173 158 158 162 169 171 167 157 163 168 ...
## $ weight   : int 65 47 43 45 58 63 47 44 48 88 ...
## $ siblings : int 1 1 2 2 1 2 2 2 2 1 ...
```

```
head(lab2, 12)
```

```
## # A tibble: 12 x 5
##       id gender height weight siblings
##   <int> <chr>   <int>   <int>     <int>
## 1    201 M        173     65         1
## 2    202 F        158     47         1
## 3    203 F        158     43         2
## 4    204 F        162     45         2
## 5    205 M        169     58         1
## 6    206 M        171     63         2
## 7    207 F        167     47         2
## 8    208 F        157     44         2
## 9    209 F        163     48         2
## 10   210 M        168     88         1
## 11   211 M        171     99         3
## 12   212 F        163     59         4
```

```
lab2test <- read.table("Lab02test.txt", header=T)
lab2merged <- merge(lab2, lab2test, by.x="id", by.y="id")
head(lab2merged, 12)
```

```
## # A tibble: 12 x 6
##       id gender height weight siblings test
##   <int> <chr>   <int>   <int>     <int> <int>
## 1    201 M        173     65         1     67
## 2    202 F        158     47         1     64
## 3    203 F        158     43         2     67
## 4    204 F        162     45         2     81
## 5    205 M        169     58         1     86
## 6    206 M        171     63         2     69
## 7    207 F        167     47         2     81
## 8    208 F        157     44         2     82
## 9    209 F        163     48         2     67
## 10   210 M        168     88         1     81
```

```
## 11    211 M      171     99      3     64
## 12    212 F      163     59      4     70
```

```
lab2merged[lab2merged$id == 211, "weight"] <- 77
lab2merged[lab2merged$height > 182, "test"]
```

```
## [1] 55 76 54
```

```
lab2f <- lab2merged[lab2merged$gender == "F",]
ordering <- order(lab2f$height, decreasing=T)
lab2f.ordered <- lab2f[ordering,]
lab2f.ordered[1, c('height', 'weight', 'test')]
```

```
## # A tibble: 1 x 3
##   height weight test
##       <int>   <dbl> <int>
## 1     178     55     47
```

5 Exercise 2.3

```
number <- 8
if (number < 100)
{ if (number < 10) { result <- "single digit" }
  else {result <- "double digits"}
} else { result <- "more than 2 digits"}

print(paste(number, "is", result))
```

```
## [1] "8 is single digit"
```

The original code is intensely messy and (to me) difficult to read. With some proper alignment and line breaks, and the modifications asked for in the question:

```
number <- 8
if (number < 10) {
  result <- "single digit"
} else if (number < 100) {
  result <- "double digits"
} else if (number < 1000) {
  result <- "three digits"
} else {
  result <- "more than 3 digits"
}
```

6 Exercise 2.4

```
speed <- 90

while (speed > 60) {
  print(paste("Your speed is", speed, "km/h! Time to slow down!"))
  speed <- speed - 7
}
```

```
## [1] "Your speed is 90 km/h! Time to slow down!"
## [1] "Your speed is 83 km/h! Time to slow down!"
## [1] "Your speed is 76 km/h! Time to slow down!"
## [1] "Your speed is 69 km/h! Time to slow down!"
## [1] "Your speed is 62 km/h! Time to slow down!"
```

```
speed <- 150
while (speed > 60) {
  if (speed > 120) {
    print(paste("Your speed is", "speed", "km/h! Slow down, BIG time!!"))
    speed <- speed - 12
  } else {
    print(paste("Your speed is", speed, "km/h! Time to slow down!"))
    speed <- speed - 7
  }
}
```

```
## [1] "Your speed is speed km/h! Slow down, BIG time!!"
## [1] "Your speed is speed km/h! Slow down, BIG time!!"
## [1] "Your speed is speed km/h! Slow down, BIG time!!"
## [1] "Your speed is 114 km/h! Time to slow down!"
## [1] "Your speed is 107 km/h! Time to slow down!"
## [1] "Your speed is 100 km/h! Time to slow down!"
## [1] "Your speed is 93 km/h! Time to slow down!"
## [1] "Your speed is 86 km/h! Time to slow down!"
## [1] "Your speed is 79 km/h! Time to slow down!"
## [1] "Your speed is 72 km/h! Time to slow down!"
## [1] "Your speed is 65 km/h! Time to slow down!"
```

7 Exercise 2.5

Here, we have a new data type: a **list**. A list in R is, as best as I can tell, similar to a dictionary/hash table in other languages.

```

statquote <- "There are three types of lies -- lies, damn lies, and
  ↪ statistics."
chars <- strsplit(statquote, split="")
class(chars)

## [1] "list"

str(chars)

## List of 1
## $ : chr [1:65] "T" "h" "e" "r" ...

chars <- chars[[1]]

```

The second example is kind of gross. It counts the number of 'e's in the string, but is unnecessarily long.

```

ectr <- 0

for (char in chars) {if (char == "e") {ectr <- ectr + 1}}
print(ectr)

## [1] 9

```

Instead, we can sum over a boolean vector, which basically counts the number of occurrences of TRUE, since T = 1 and F = 0.

```

sum(chars == "e")

## [1] 9

sctr <- 0

for (char in chars) {
  if (char == "s") sctr <- sctr + 1
  if (char == "m") break
}
print(sctr)

## [1] 3

```

8 Exercise 2.6

```
store1 <- c(15, 9, 13, 5, 2, 17, 14)
store2 <- c(16, 7, 5, 16, 8, 13, 14)
```

```
store3 <- store1 + store2
store3
```

```
## [1] 31 16 18 21 10 30 28
```

```
mean(store3)
```

```
## [1] 22
```

```
store3 <- store3[order(store3)]
store3
```

```
## [1] 10 16 18 21 28 30 31
```

```
mean(store3, trim=0.2)
```

```
## [1] 22.6
```

The optional `trim` parameter of `mean()` takes out a percentage of the highest and lowest values before calculating the mean. For example, if `trim=0.1`, then the highest 10% and lowest 10% of values are removed.

The number of values trimmed is rounded down. For example, taking the mean of the above vector of 7 items:

- if `trim = 0.1`, trim 10% from each side, or 0.7 values, rounded down to 0 values. No trimming occurs.
- if `trim = 0.2`, trim 20% from each side, or 1.4 values, rounded down to 1 value.
- if `trim = 0.4`, trim 40% from each side, or 2.8 values, rounded down to 2 values.

9 Exercise 2.7

```
m2 <- function(x) {
  y = (x - mean(x))^2
  return (mean(y))
}
```

```
m4 <- function(x) {  
  mean((x - mean(x))^4)  
}
```